International Journal of Computer Techniques – Volume 12 Issue 1, January - Feb - 2025

AI Driven ChatOps with for Realtime Security Incident Response in DevSecOps

Balajee Asish Brahmandam
University of Texas at Austin
Austin, Texas, USA
balajeeasish@utexas.edu

***Abstract*— DevSecOps techniques speed up software delivery, while security incident response occasionally lags in speed and integration. This study provides an artificial intelligence-powered ChatOps framework to streamline and automate security incident response inside DevSecOps procedures. We demonstrate how ChatOps (chat-centric operations) and artificial intelligence (AI) enable real-time automatic remediation, cooperative investigation, and quick detection of security issues. The technique integrates machine learning-based threat detection with team chat platforms (like Microsoft Teams and Slack), allowing security bots to use chat instructions to perform containment operations, assist in triage, and notify teams of incidents with more information. We evaluate the benefits of this integration, such as improved teamwork, less human error, faster reaction times, and continuous event learning. In a case study of the banking sector, automated documentation demonstrates a significant improvement in compliance and reaction speed (minutes as opposed to hours). We discuss topics like automation confidence, false positives, and tool integration, and we suggest future advancements like advanced language models for incident handling. The results demonstrate how ChatOps, powered by AI, might transform incident response into a rapid, proactive process that aligns with DevSecOps agility.**

Keywords*— *DevSecOps, ChatOps, AI in Cybersecurity, Security Automation, Incident Response, Security Orchestration Automation and Response (SOAR), Machine Learning, Security Bots, Real-Time Threat Detection, Natural Language Processing, Collaboration Tools, Continuous Monitoring, SIEM, Microsoft Sentinel, IBM QRadar, GPT-powered Assistants, Financial Technology, Cyber Threat Intelligence.

I. Introduction

Modern businesses have embraced DevSecOps to integrate security into the rapid DevOps lifecycle, even though incident response processes remain weak. Traditional incident response is frequently reactive, delayed, and isolated from development, making it challenging to keep up with the rapid pace of continuous deployment. Alert fatigue brought on by a high volume of security alerts causes practical missed or delayed reactions from security professionals. For example, the security department of one bank claimed that the time-consuming nature of human examination of each warning resulted in delayed reactions and increased exposure risk. In a time when cyberattacks are becoming more frequent, delays in discovery and containment can be costly. NIST SP 800-61 and other industry standards stress that "rapidly detecting incidents" depends on an effective incident management capability and aids in damage mitigation. Therefore, incident response needs to be as fast and flexible as DevOps.

ChatOps has emerged as a promising model for bridging this gap by combining operations and tools into a collaborative dialogue area. Team communication channels serve as the interface for running tools and scripts in the "chat operations" (or "ChatOps") style of working. According to one practitioner, "ChatOps is a collaboration model that connects people, tools, process, and automation into a transparent workflow," allowing conversations to be immediately turned into action.

In a ChatOps-driven incident response, all relevant warnings, information, and corrective actions are routed through an ongoing chat room. Because everyone can see the same information and context in real time, this approach tightens feedback loops for decision-making and breaks down silos. Teams no longer have to switch between consoles, emails, and ticketing systems during a crisis; the chat serves as the event management war room. However, centralizing communication alone is not enough when threats are changing at the speed of machines. The volume and intricacy of data may still be too much for human responders to handle.

This is where integrating artificial intelligence is essential. Machine learning and security analytics advancements make it possible to identify threats and anomalies more quickly than human analysts can. AI-driven security solutions, such as Microsoft's cloud AI analytics or IBM's Watson-powered security intelligence, may automatically search through records and find questionable patterns. We can go from a purely reactive approach to a more proactive and automated incident response by integrating ChatOps with artificial intelligence. Intelligent bots in an AI-powered ChatOps system continuously scan systems, provide contextually rich alarms, and even initiate containment measures before a human is involved. People can then verify, discuss, and direct the artificial intelligence's actions through the cooperative dialogue. This mix of human oversight and machine speed promises to cut response times from hours to minutes and

better align security response with DevSecOps' continuous delivery mindset.

This paper outlines a comprehensive approach to AI-driven ChatOps in DevSecOps and assesses how well it automates security incident response. Section II examines enabling technologies and related activities and includes significant industry solutions and recommendations. Section III discusses the recommended architecture and process for integrating ChatOps and artificial intelligence. Section IV discusses the outcomes and benefits and includes a financial sector case study demonstrating the real impact. In Section V, we examine potential future advances, such as more advanced artificial intelligence assistants, and critically evaluate the challenges and limitations of AI-driven ChatOps. Finally, Section VI reveals that AI-powered ChatOps may significantly increase incident response speed and collaboration, fortifying businesses against assaults without sacrificing DevOps agility.

II. LITERATURE SURVEY

Both business and standards have long recognized the need for quick, automated incident response in DevSecOps. The NIST Computer Security Incident Handling Guide (SP 800-61) emphasizes that businesses must respond quickly to reduce harm and strongly emphasizes timely discovery and reaction. Security operations centers rely on playbooks or runbooks for response, and Security Information and Event Management (SIEM) systems for warning collection in traditional setups. However, these could be slow if they are not part of the development process. To include security in rapid DevOps operations, DevSecOps requires continuous and smooth incident response. This has increased interest in techniques that automate incident management and include developers.

AI-assisted collaboration in incident management has been the focus of recent research. Brooks et al. (2025) examine how AI-powered ChatOps might improve incident resolution and collaboration efficiency in DevOps teams, reflecting growing scholarly interest in this convergence of artificial intelligence, DevOps, and SecOps. The idea of ChatOps, which uses chat platforms to speed up operations, gained popularity in the middle of the 2010s, particularly thanks to GitHub and others. The advantages of ChatOps in incident management have been documented by Atlassian and other DevOps tool manufacturers, who emphasize how it keeps everyone in sync and breaks down information silos.

That channel creates a consistent event timeline by recording all discussions and activities. By centralizing communication, ChatOps encourages transparency and sharing situational knowledge during incidents. Early ChatOps adopters in operations have reported improved team culture and faster coordination. "Security ChatOps" is a conversation-driven inquiry in security, where analysts, tools, and bots collaborate within a single window. This concept has given rise to plugins and solutions that connect chat platforms to security systems, allowing analysts to do incident response tasks and pose queries directly from the dialogue. For example, the Palo Alto Networks security team noted that ChatOps might improve transparency because all analysts can see the commands carried out and potentially speed up response times by eliminating the need to switch between multiple interfaces. These market insights lay the foundation for integrating artificial intelligence (AI) into the ChatOps loop.

Security Orchestration, Automation, and Response (SOAR) technologies have grown significantly in tandem with ChatOps. By codifying incident response procedures, SOAR platforms like IBM Resilient, Splunk Phantom, and Palo Alto Cortex XSOAR enable security teams to automate playbooks. Once specific requirements are met, these systems can automatically perform tasks like alert enrichment, notice delivery, or endpoint quarantine. Artificial intelligence integrated into security monitoring has further enhanced these functions. IBM QRadar Advisor is a noteworthy example. IBM's QRadar SIEM integrates Watson artificial intelligence to analyze threats and provide insights that typically need a skilled analyst.

Watson can automatically investigate an incident and produce a report in natural language with findings and confidence scores based on threat intelligence and historical data. ChatOps can then be used to implement such AI-driven analysis. IBM has demonstrated ChatOps interfaces in which Watson's suggestions, such as suggesting a solution or identifying a potential root cause, are shared in a Slack channel for the analysts to review. This demonstrates how AI could be a force multiplier by handling the heavy cognitive strain of analysis and freeing people to focus on making decisions.

Cloud firms have corporate artificial intelligence in their security offerings. Microsoft Sentinel (previously Azure Sentinel) is a cloud-native SIEM that uses integrated machine learning and User and Entity Behavioral Analytics (UEBA) to identify threats and anomalies.

Sentinel's tight integration with automation procedures is one of its key benefits; it makes it possible for Azure Logic Apps playbooks to initiate incident responses. According to Microsoft documents, there are instances where a Sentinel rule that detects a high-priority incident will automatically alert the on-call analysts by sending a message to a Microsoft Teams channel. Analysts can then use Teams to collaborate and even approve some automated tasks. This implements a ChatOps model: Sentinel, the AI brain, detects an issue and notifies users through a chat interface (Teams), allowing them to direct the subsequent course of action. Similar interfaces are available from other vendors; Palo Alto Cortex XSOAR, for

instance, provides pre-built Slack and Teams bots that mimic its incident "war room."

Analysts may investigate and manage events directly from Slack since all operations are linked to the central SOAR console. This suggests the Slack bot could acquire firewall logs or do a virus scan, and the results would be sent over chat. These integrations demonstrate that ChatOps offers a unified interface for managing complex, multi-system interactions that use automation and artificial intelligence in the background and are convenient.

III. METHODOLOGIES

Our proposed method, AI-Driven ChatOps for Incident Response, integrates automated action execution, collaborative communication, and threat identification based on artificial intelligence.

We outline the technical integration strategy, stage by stage, along with the enabling elements.

1. AI-enabled security technologies begin the pipeline with continuous monitoring. AI-enabled security products begin the pipeline with continuous monitoring. This includes SIEM systems with machine learning analytics, intrusion detection systems, endpoint detection and response (EDR) technologies, and more. These tools are designed to identify trends of known threats or anomalies. For example, Microsoft Sentinel may use anomaly detection models to monitor network traffic and logins, whereas an EDR may use behavioral analysis to detect malware. The system immediately generates an incident or alert when a suspicious occurrence or pattern is identified confidently. Crucially, the AI/analytics system alerts the DevSecOps team's ChatOps channel instead of merely documenting the event. Integration hooks do this: for example, when an incident is produced, a Sentinel playbook or custom script sends a message to a Slack/Teams channel using webhooks or bot APIs. The chat alert typically includes a synopsis of the issue and pertinent background data. Automating the detection-to-notification process eliminates any lag between a security event occurring and the team learning about it.
2. The system provides further context to aid in immediate comprehension once an alert emerges in the ChatOps channel, such as a Microsoft Teams issue conversation or a specific "#security-incidents" Slack room. The ChatOps bot, an integrated notification system, includes important data gathered by artificial intelligence and the warning (e.g., "Possible data exfiltration detected on Server X"). A timestamp, the affected host or application, the triggering anomaly (such as "outbound data spike of 500MB to unknown IP"), and a risk or priority score, for instance, could all be included in the message. By connecting with knowledge bases and past incident data, the warning can be linked to similar occurrences, for example, "Pattern resembles incident #248 (SQL injection attempt) last month." At this stage, the bot might even make recommendations in some versions. This is demonstrated by IBM's Watson-based solutions, which automatically provide a suggested diagnosis or course of action. For example, Watson may determine that there has been "suspected credential compromise" and suggest "reset user password and investigate login origin." To ensure that on-call responses are aware of the warning immediately, the chat system alerts all channel team members, whether they are mentioned or not (typically via mobile push notifications). Even later joiners can scroll up to see the entire alert's history and background, thanks to an ongoing chat.
3. Since the team can now see the incident, the following step is triage, which artificial intelligence speeds up. AI-assisted investigation and triage: Since the team can now see the issue, the next step is to expedite the triage process using AI's assistance. Team members can interact with the ChatOps bot to request additional information or perform quick checks, all in plain English or by following instructions. If an analyst were to say, "Bot, show related alerts in the last 24h for Server X," for example, the SIEM-integrated bot would be able to retrieve and present the relevant information rapidly. This eliminates the need for the analyst to search many systems manually.

Machine learning comes into play when providing intelligent assistance. The bot can be set up to understand common queries and obtain answers through Natural Language Processing (NLP), acting as a security chatbot. Furthermore, the artificial intelligence can prioritize the occurrence by connecting it to other signals. The artificial intelligence can highlight the situation as a single high-priority event if many low-level alarms are coupled. Artificial intelligence now assists less experienced team members by recommending likely causes or subsequent steps based on historical patterns. The bot might respond, for instance, "This alert is probably a false positive, similar to events seen before, because X, Y, Z," or even "This seems serious; 85% probability of ransomware infection affecting three hosts".

Such guidance is provided by trained models that use past incidence data. With the proper focus, the team can quickly understand the scope and impact of this AI-powered triage.

1. ChatOps excels in this stage because it allows the chat interface to act quickly while the human

operators, who may include developers, operations engineers, and security analysts, use the chat as a collaborative tool to decide on containment or corrective actions based on the information available. The team can issue commands to the bot by interacting with it through interfaces with orchestration tools like automation scripts or SOAR platforms. For instance, the channel may provide the following interactive prompt: "Quarantine Server X? When a team member clicks "Yes" or gives an approval command, the bot will initiate a predetermined containment plan in response to the interactive prompt on the channel: "Quarantine Server X? (Yes/No)." Chat instructions can initiate actions like patch deployment, service restart, user account deactivation, host isolation from the network, and message with an automation engine connected to a Teams bot. For instance, the bot publishes "Executing containment: isolating DBServer12 from the network" so that all participants in the chat can immediately see the action being taken. Even when one team member takes the initiative, ChatOps' collaborative nature enables multiple team members to contribute and discuss simultaneously. The virtual "war room" is the conversation; calling another phone is unnecessary. This guarantees everyone knows the most recent scenario and quickens decision-making. The AI still assists during reaction; for instance, if the team is unsure about a course of action, they can ask the bot for a risk assessment, and the AI could respond with likely outcomes. Combined with human judgment and artificial intelligence speed, routine procedures are automated under human supervision for critical decisions. This method is sometimes called "human-in-the-loop automation."

The ChatOps bot updates the channel on any automatic tasks while the issue is being handled. If the bot acts, like blocking an IP at the firewall, it reports the result (success or failure and any pertinent details. The chat contains the entire history of events, including the initial alarm, team discussions, instructions conveyed, and the final resolution. This transcript becomes a valuable artefact after the immediate threat has been managed. It can be automatically exported to a ticket or issue management system (like Jira or ServiceNow) for documentation. Many compliance regulations require a post-incident report; in this case, the chat log provides a precise chronology of who did what and when.

Additionally, the AI systems use this data to improve their subsequent performance. Machine learning models can be trained using the event's characteristics and outcomes (with appropriate data handling). If the AI makes a helpful suggestion, it will learn more; if it misses an indication that a human saw, that information might be incorporated into the logic for further detections. The result is continuous improvement: the AI becomes increasingly accurate in identifying and diagnosing issues over time, and the playbooks may be refined based on the team's findings. The ChatOps approach inherently promotes post-incident assessment since all decisions and insights are logged in one place. The conversation can be used for team retrospectives to identify misunderstandings or delays. This closed-loop learning is a feedback loop between artificial intelligence and human expertise, where each event handled by AI-ChatOps increases the system's intelligence (via machine learning training) and the team's knowledge (through accumulated lessons).

Implementation Architecture: This system is made up of several parts: (a) security analytics driven by AI that feature alerts and API capabilities, like a SIEM/XDR platform or custom ML models for anomaly detection; (b) a ChatOps platform that the team frequently use, such as Microsoft Teams, Slack, or another comparable tool; (c) an integration layer or chatbot that links the security tools to the conversation (this might be a bot application that posts messages and listens to alarms from the AI tool, as well as one that watches for commands in chat to carry out actions using the security tools); and (d) an automation/orchestration engine that the chatbot uses to carry out response activities (such as a script runner or SOAR platform). Communication typically relies on secure webhooks or APIs. For example, Microsoft Teams provides an incoming webhook URL to which a script can write, whereas Slack offers event APIs and bot tokens for two-way communication. When an alert is generated, the AI analytics system is set up to initiate a webhook or send a message, for instance, by utilizing a Python script in QRadar's case or a Logic App in Azure. On the other side, the chatbot is configured with handlers for certain chat buttons or phrases; The bot receives the command when a user writes @SecBot isolate hostX and initiates the appropriate automation workflow, which could be a confinement script, an API request to a firewall, etc. We ensured that permissions and authentication were followed at every turn to prevent illicit activity. For instance, the bot can have a restricted allow-list of commands for non-admin users or require a second confirmation.

Our system essentially coordinates these integrations into a standardized incident response procedure. The results of implementing such an AI-driven ChatOps strategy are provided in the next section, along with a detailed scenario demonstrating its effectiveness and the benefits observed.

IV. RESULTS

After the text edit has been completed, the paper is ready for the template. Duplicate the template file by using the Save As command, and use the naming convention prescribed by your conference for the name of your paper. In this newly created file, highlight all of the contents and import your

prepared text file. You are now ready to style your paper; use the scroll down window on the left of the MS Word Formatting toolbar.

Businesses would realize significant improvements in overall operational efficiency and significant incident response statistics by implementing the AI-driven ChatOps methodology. The main advantages noted align with the principles of DevSecOps, which include speed, teamwork, and continuous development. Our research and industry studies serve as the basis for the following findings:

Faster Reaction Times: The time between detection and initial response is significantly reduced by automation. Responders immediately receive notifications from chat, and within seconds of detection, containment actions (such as isolating a compromised server) can be initiated. Mean Time To Respond (MTTR) is decreased as a result. Automated playbooks might complete critical tasks before human team members mobilize. Compared to traditional workflows, where an event may go unnoticed for minutes or hours and require manual coordination, AI-ChatOps can drastically reduce response time. Rapid confinement reduces the time an attacker can stay, potentially reducing damage. This speed is critical in industries like finance, where even a few minutes of unchecked malicious activity can have significant consequences.

By consolidating all incident data and communications into a single channel, ChatOps ensures that developers, operations, and security staff have the same situational awareness. All participants have real-time access to the data, actions, and warnings; No separate segregated discussions or parallel email channels exist. This transparency eliminates the "telephone game" of misunderstandings and allows decisions based on a common operating picture. AI facilitates team decision-making by instantly incorporating expert knowledge, such as threat intelligence context or recommended actions, into the conversation. This improves the team's feedback loop and fosters a collaborative culture where responding to incidents is a shared duty rather than solely the task of a separate security team. AI-driven ChatOps automates time-consuming and stressful tasks, reducing the likelihood of mistakes that typically occur when people are rushed or exhausted. For example, carrying out a challenging confinement procedure at three in the morning.

Using a tried script rather than manually typing a series of commands into a console is more reliable. The AI can also serve as a safety net before execution by verifying commands and alerting users if an activity appears unusual or potentially dangerous. Since all tasks are completed through a single interface—the chatbot—playbooks handle the known steps consistently, lowering the possibility of forgetting a step or misconfiguring a tool. Errors brought on by misunderstandings are also avoided because there is just one source of truth in the conversation log; there is never any doubt about what was done or who is doing what.

Each event that is controlled through the ChatOps channel becomes a lecture recording. By reviewing the chat transcript, teams may conduct post-mortems and determine what worked and could be improved. Future playbook modifications and enhanced techniques follow from this. From the artificial intelligence perspective, machine learning models can be fed event data, such as warnings, actions taken, and outcomes. Over time, artificial intelligence improves its ability to distinguish between benign patterns and those that point to issues. For instance, if the AI detected an event that later proved false positive and the team flagged it, the model may learn from that information. However, such information could help with future recommendations if the team disregarded an AI tip in favour of a different action that resolved the problem.

This positive feedback loop suggests that the system may become more effective the more it is used. New team members can quickly increase their understanding of the organization's security situation by reading through the archives and learning from past incidents. AI-powered ChatOps integrates security monitoring and response directly into CI/CD workflows by merging them with tools and procedures that DevOps teams were already familiar with. IT professionals and developers frequently use chat platforms for operations and deployments, such as executing chat commands to deploy code or check system status. By using the same interface for security events, handling them becomes an organic part of everyday tasks rather than a separate endeavor.

Developers are encouraged to be more involved with security issues since they are evident in a forum where they participate. Additionally, ChatOps can communicate with CI/CD pipeline events. For example, if a significant vulnerability is found in a new build, the system might raise an alarm during the chat and even halt the pipeline until it is corrected. Enforcing security gates in real time through chat ensures that security is integrated into the development process. By interacting with DevOps teams in chat, ChatOps with AI contributes to the DevSecOps goal of integrating "security as code" and collaboration rather than viewing security as a separate silo or an afterthought.

To demonstrate these benefits, we use a case study in the financial services industry, where regulatory constraints and high risk make incident response accuracy and timeliness top concerns. This scenario is a combination of typical threat events in a bank and the real capabilities of AI-SecOps tools: Consider an AI-enabled security monitoring system that is integrated into the IT architecture of a large retail bank. The system consists of either Microsoft Sentinel or IBM QRadar with Watson, SIEMs that use logs from servers, databases, and network devices. The AI analytics engine discovers an

anomaly late one evening: DBServer12, a normally inactive database server, is suddenly sending a huge volume of data to an unknown external IP address.

This pattern points to a possible breach in progress—an attempt to exfiltrate data. This is immediately recognized as a high-priority occurrence by the artificial intelligence. It determines a significant chance of malicious activity by correlating several indicators, including the amount of data, the process initiating the transfer, and comparisons to established attack patterns. Instead of waiting for a human analyst to notice this in a dashboard, the system initiates an incident and sets out a ChatOps alert. The bank's "#security-war-room" Slack channel receives the following message a few seconds later: "Alert: Possible data exfiltration detected on DBServer12". The message provides an overview: "DBServer12 sent 900MB of data to IP 203.0.113.45 (Russia) at 23:10 UTC, exceeding baseline by 50x."

Additionally, it provides background data, such as a hazard score and the username carrying out the process. "This pattern resembles the ACME Corp breach (June 2024) in our knowledge base," the machine would say, based on its past knowledge. The bot automatically tags the on-call security engineer and relevant team members (network engineer, database administrator, etc.) to ensure their devices receive an alert notification immediately.

Watson, the AI in this case, gives a recommended course of action along with the warning: "Suspected data theft in progress. In light of earlier responses that successfully stopped similar attacks, I suggest turning off database user ' svc_import' and stopping DBServer12's outgoing connectivity." Currently, the team meets in the Slack channel; some members use computers, while others use phones for remote work.

Everyone is looking at the same data the bot posted in the first minute. "This looks like a probable insider or malware exfiltration; let us contain first," the security engineer says as the cooperative discussion begins. "That server holds customer data; we must ensure backups are intact too," the database administrator says. These conversations are taking place in a single thread accessible to all parties involved, including the compliance and communications team members who joined in response to the high-priority alert, thanks to ChatOps.

Without wasting any time, the crew proceeds with the recommended containment. Through the ChatOps interface, the security engineer inputs a command to the bot: Block-outbound DBServer12 @SecBot.

The bot, which is integrated with the bank's network firewall controller (also known as the SOAR platform), isolates that server by running a script that cuts off its external network access. Another command, @SecBot disable-user svc_import, runs a script that removes the suspicious user's directory service credentials. The potentially dangerous data transfer is halted within one to two minutes following the initial notification. Compare this to a traditional method without ChatOps/AI, where it may have taken the team 15 to 30 minutes to figure out what was going on, find everyone, join a call, and connect to multiple systems to limit the server. Here, critical containment occurred very quickly, greatly minimizing any data loss.

The Slack channel logs every step the bot does, including "Firewall rule added: blocking DBServer12 outgoing traffic" and "User svc_import locked." These activities are observed by the compliance officer and IT operations staff, who are also in the channel. Compliance claims that all activity is timestamped and recorded, which will be helpful for reporting. Financial regulations sometimes require a timeline of event activities; the conversation log often provides that on the spot. The result of the bot indicates whether any memory-found malware was present. In the meantime, the communications manager drafts a customer notification form if this becomes a breach requiring disclosure. No emails are sent; this is done over the chat operations channel.

Within fifteen minutes of detection, the immediate threat is eliminated. The data exfiltration was stopped, and the used account was deleted. A further investigation reveals that an employee's account was compromised and used to export illegal data. Due to the prompt response, the damage was minimal. However, the bank does adhere to strict breach response procedures; organizations typically have 24 to 72 hours to notify authorities of any significant incident; in this case, they received a comprehensive report in less than an hour. The Slack incident channel's record is kept in the incident ticket.

The security team and the AI will retrain detection models using the data from this incident in the following days, taking note of the indicators (so that if a similar pattern emerges, it might be discovered even sooner). This scenario highlights a few practical advantages of ChatOps powered by AI. The prompt identification and response prevented a significant data compromise. Everyone involved was informed because of the collaboration and transparency, and nothing was overlooked or delayed because of a single person. The team felt empowered to move quickly because the AI had done some of the analysis that humans would generally take longer to accomplish.

The event was also easily recorded, which is crucial for post-event evaluation and financial compliance. It is similar to a real-world case study where a bank experiencing alert fatigue and delayed, manual responses improved their performance by utilizing integrated tools and automation. Before automation, the security staff in that real-world situation (NKGSB Bank)

was overworked and occasionally responded too slowly. However, they significantly reduced alert quantities and reaction times after implementing an integrated system. Our scenario demonstrates the same idea: financial institutions (and businesses in general) may stay ahead of risks while maintaining the audit trails and confidence that stakeholders and regulators require by integrating AI with ChatOps.

V. DISCUSSION

Despite the promising results of AI-driven ChatOps for incident response, there are still important problems and things to work on. There are difficulties in using this approach, and certain limitations must be tolerated. We also look into potential future directions that could further enhance AI-ChatOps' effectiveness in DevSecOps.

A. Problems and Limitations: One major challenge is integration complexity. Because enterprises typically have a diverse mix of security solutions (SIEM, EDR, firewalls, ticketing systems, etc.), it takes a lot of engineering work to get all these to work together seamlessly with a chat platform. Some tools do not have built-in chat connections or APIs. Teams may need to use middleware like a SOAR platform or write custom scripts to fill in any gaps.

This integration effort could be significant, especially in organizations with outdated systems. Additionally, ChatOps introduces a learning curve and a culture shift. Team members must get used to carrying out important duties in a chat interface. At first, some people may be skeptical or concerned about a bot changing systems. Gaining trust in the automation is essential, which typically entails extensive playbook testing and perhaps starting with suggest-mode or read-only bots before allowing them to alter systems. There is a risk associated with an over-reliance on AI. If team members begin to mindlessly follow AI guidance, a mistake in the AI's logic could lead to incorrect behaviours. A false positive warning, for example, may result in an unnecessary server shutdown.

To counteract this, human validation—the person-in-the-loop concept—should continue to be used for important activities. The AI should be prevented from behaving outside predetermined bounds by stringent access restrictions. The fact that current artificial intelligence models are not perfect despite their advancements is another limitation. Attackers may devise ways to evade machine learning detection or even manipulate artificial intelligence by triggering many noise alarms to distract the AI or providing adversarial inputs. The system must be designed to handle false positives and negatives elegantly. False negatives, or missed events, are a concern if the team becomes unduly dependent on the AI to produce alerts.

Therefore, traditional monitoring and redundant detection methods should be used as backups. False positives in the chat can result in "alert spam," which could cause severe alerts to go ignored (the classic boy-who-cried-wolf problem). It takes constant work to fine-tune the AI models and alert thresholds; feedback loops are necessary in the AI-ChatOps system to continuously increase accuracy. Scalability and performance must also be considered. For example, a large firm may generate thousands of daily notifications, which would overwhelm users if sent to a chat channel. The issue is resolved by letting the AI filter and summarize, but again, this requires very high artificial intelligence performance.

The security of the ChatOps channel itself is another crucial subject of discussion. Ironically, if the chat platform is not adequately secured, it could serve as a conduit for attackers. Private incident data could be viewed, or potentially destructive instructions could be sent by an adversary with access to the incident channel or the bot's credentials. Thus, encryption and access controls are necessary for secure tokens for bots and strong multi-factor authentication for human users. The bot should be handled similarly to a privileged account; its actions should be monitored, and its login credentials should be secured. There should be fail-safes as well. For instance, some hazardous operations would require dual clearance via chat (two authorized people checking) to be carried out. This would prevent a scenario in which damage may be caused by either an AI error or a single compromised individual.

There is a change management component from the perspective of individuals. Teams need training to maximize the use of the newest technologies. Some traditional incident responders may object to the chat-centric approach and would rather work over phone bridges or on well-known SIEM consoles. It might be helpful to demonstrate the benefits and provide gradual adoption, starting with ChatOps for less significant events and expanding from there. It is also critical to establish protocols for the use of the conversation, including etiquette, responsibilities, how to escalate outside of chat if necessary, etc.

It is crucial to ensure that the ChatOps approach complies with corporate or regulatory regulations, such as data retention and privacy, as chat logs may contain sensitive information that needs to be handled per the policy.

B. Prospective Advancements: Despite these challenges, technology's trajectory suggests that many issues may be mitigated, and opportunities will only improve. One important area for future development is using advanced artificial intelligence, especially generative models, in the ChatOps process. This is already beginning: complex queries may be analyzed by large language models (LLMs) like GPT-4, which can provide explanations like a human.

A future ChatOps bot might use an LLM in incident response to examine unstructured logs or provide more contextually

aware response suggestions. For example, the bot might converse with analysts rather than providing pre-programmed responses. For example, an analyst can ask, "What is the possible impact of this incident?" An AI like GPT might combine an answer with event data and prior knowledge. "What is the possible impact of this incident?" is a plain-language query that an AI such as GPT may combine with an answer derived from incident data and existing knowledge.

Microsoft's launch of Security Copilot, an LLM-based AI assistant for cybersecurity, is a start in this direction. Its purpose is to assist analysts by summarizing events and suggesting questions that might easily link with a ChatOps interface in the future. A similar future where regular paperwork is fully automated by reading the natural language conversation is hinted at by Atlassian's recent AI tools that condense chat conversations into Jira tickets.

Another expected development is greater cross-border integration. Nowadays, ChatOps is primarily limited to a single team or business environment. In the future, we may have secure cross-organizational threat intelligence or plan reactions to significant occurrences, such as critical vulnerabilities affecting numerous companies.

These conversations may be mediated by AI, which would instantly pull relevant data from external sources. Enhancements in standardization, such the development of standard formats for incident data and response playbooks, may also facilitate the collaboration of different technologies in an AI-ChatOps system. The OASIS CACAO standard for cybersecurity playbooks and the STIX/TAXII standards for threat intelligence are two examples. This suggests that a business might install a new detection technology without requiring much bespoke programming and have it start feeding the chatbot.

Additionally, methods for human-AI collaboration will advance. Instead of AI only setting off alarms and people responding, future systems might be conversational.

To make the conversation more interesting, the AI might, for instance, ask the team clarifying questions: "This appears to be a potential false positive. In your opinion, should I disregard such alerts for the following hour?" The concept of conversational AI for security can significantly increase productivity since it may be able to make more decisions when instructed. We anticipate that chat platforms' user interfaces will evolve to suit these workflows better, such as using mini-forms for approval or richer message formats with integrated dashboards (some already present in platforms like Slack and Teams, but could be more AI-driven).

In order to manage current limits, explainable artificial intelligence (XAI) research will be essential. Suppose the AI can explain why it made a recommendation—for instance, by emphasizing the factors that led it to flag an incident as high priority—human operators will be more likely to trust and comprehend it. If necessary, they will also be able to fix it. This might be included in the chat conversation; the bot could explain any alarm it posts if the user requests it.

From the standpoint of organizational processes, incident response procedures may eventually be reorganized to formally integrate AI and chat when AI-driven ChatOps proves its value. Using collaborative technologies, laws, and cybersecurity frameworks may evolve to acknowledge controls and record-keeping.

Companies may start monitoring MTTR and other metrics, especially the speed at which the ChatOps loop handles an incident and use that information to drive improvements.

However, it is essential to consider the human element carefully. Enhancing human decision-making, not replacing it, is the goal. Our discussion highlights that, despite automation, human oversight is necessary to handle novel or complex situations that AI might not fully comprehend. There will always be incidents that defy automation, such as multifaceted attacks that need creative investigation or zero-day attacks that lack a known pattern. The ChatOps architecture is still helpful as a coordinating tool for these individuals, even if the AI withdraws.

While problems like accuracy, trust, and integration effort need to be managed, incident response will increasingly use artificial intelligence in collaborative settings.

As technology and methods advance, AI-driven ChatOps may eventually become a standard part of DevSecOps toolchains, much like CI/CD pipelines and infrastructure-as-code are now. Those that stick to manual and segregated processes may find it more challenging to keep up with the speed of software delivery and cyber risks; early adopters will benefit from resilience through speed and shared intelligence.

VI. CONCLUSION

AI-driven ChatOps significantly advances how businesses handle security events in a DevSecOps environment. In this paper, we reconstructed the idea of integrating ChatOps and artificial intelligence, demonstrating how it combines the best aspects of both worlds: human creativity and teamwork in problem-solving with computers' speed and analytical power.

By integrating artificial intelligence capabilities into chat-based procedures, DevSecOps teams may detect incidents sooner, react faster, and collaborate more effectively without

leaving their regular communication channel. Continuous delivery rules transform the incident response process from a disorganized, occasionally frantic rush into a clear, consistent strategy.

Our analysis of current technologies and literature showed strong support for this integrated approach, ranging from NIST's calls for prompt action to industry case studies that highlighted the shortcomings of purely manual processes to successes with automated, collaborative solutions. We described how artificial intelligence monitors systems and sends comprehensive notifications to a ChatOps channel, where the team and AI collaborate to identify and neutralize the danger. The results are significantly quicker reaction times, improved communication between the development and operations teams, and event data gathering for continuous improvement. With technologies like IBM QRadar, Microsoft Sentinel, and Palo Alto XSOAR under consideration, the banking industry example demonstrated how AI and ChatOps working together may stop a potentially disastrous attack in minutes.

Critical emphasis was paid to the challenges of implementing AI-driven ChatOps, such as integration problems and the need to maintain human oversight. Careful planning, training, and iterative system improvements may overcome these fundamental challenges. Furthermore, the benefits outweigh the drawbacks in situations where every second counts during an event.

Additionally, we observed that this approach fosters a culture in which security is "baked in"—that is, engineers and analysts collaborate in a single forum, sharing responsibilities rather than segregating functions. This shift in mindset aligns with DevSecOps principles and could lead to more secure development practices in general. We anticipate that ChatOps powered by AI will become increasingly commonplace and maybe necessary. Manual responses can no longer keep up with the increasing complexity of cyberattacks. However, as software delivery speeds up, businesses seek security solutions that do not impede innovation. Both goals are met by AI-driven ChatOps, improving security responsiveness without taking engineers off their workflow.

These chat assistants could become much more intelligent and practical thanks to advancements in artificial intelligence, especially in knowledge representation and natural language interpretation. An AI assistant has already finished the routine tasks and provides the team with clear resolution options, so it is reasonable to assume that a security warning will soon be less of a cause for concern and more of a staged conversation.

Ultimately, ChatOps and artificial intelligence work well together to automate incident response's repetitive and time-consuming tasks while preserving human control over crucial choices. Businesses that use this approach can improve their security posture, lessen the impact of incidents, and maintain the rapid pace of DevOps delivery.

Research and case studies indicate that AI-driven ChatOps is a game-changing method for DevSecOps rather than merely a minor enhancement. In an era of constant change, as threats evolve and our tools improve, such collaborative automation will likely be a key component of maintaining safe and robust systems.

REFERENCES

[1] NIST, *Computer Security Incident Handling Guide (SP 800-61 Rev.2)*, National Institute of Standards and Technology, 2012.

[2] S. Regan, "ChatOps is a collaboration model that connects people, tools, process, and automation...," *Atlassian* (Incident Management blog), 2020.

[3] Palo Alto Networks, "Faster incident response and reduced alert fatigue at NKGSB Bank," Case Study, Palo Alto Networks, 2023.

[4] Microsoft, "Use a Microsoft Sentinel playbook to stop potentially compromised users," Microsoft Learn (Tutorial Documentation), 2023.

[5] IBM Security, "IBM QRadar Advisor with Watson," IBM Product Brief, 2017.

[6] J. Goh, "Frequently Asked Questions About Security ChatOps," *Palo Alto Networks Blog*, Nov. 2017.

[7] M. Du, F. Li, G. Zheng, and V. Srikumar, "DeepLog: Anomaly Detection and Diagnosis from System Logs through Deep Learning," in *Proceedings of the 2017 ACM Conference on Computer and Communications Security (CCS)*, 2017, pp. 1285–1298. https://doi.org/10.1145/3133956.3134015

[8] A. Gember-Jacobson, R. Viswanathan, and S. Akella, "OpenSOC: A Security Operations Center Framework for AI-Augmented Monitoring," *IEEE Security and Privacy*, vol. 18, no. 5, pp. 25–34, 2020. https://doi.org/10.1109/MSEC.2020.3005083

[9] A. Newell, B. Aragam, and C. C. Aggarwal, "Harnessing LLMs for Security Operations: A Framework for Natural Language-Based Alert Triage," *arXiv preprint arXiv:2302.09831*, 2023. https://arxiv.org/abs/2302.09831

[10] R. Ross et al., "Security and Privacy Controls for Information Systems and Organizations," *NIST Special Publication 800-53 Rev. 5*, U.S. Department of Commerce, 2020. https://doi.org/10.6028/NIST.SP.800-53r5

[11] Google Cloud, "Chronicle SOAR: Threat Investigation and Response Made Easy," *Whitepaper*, 2023. https://cloud.google.com/chronicle/docs/soar

[12] GitHub, "ChatOps: Conversation-driven Development," *GitHub Blog*, 2021. https://github.blog/2021-06-10-chatops-conversation-driven-development/

[13] IBM, "Security and Resilience with Watson for Cybersecurity," *IBM Research White Paper*, 2020. https://www.ibm.com/security/watson

[14] M. Miller and J. Osborn, "DevSecOps Essentials: Security as Code in the CI/CD Pipeline," *IEEE Software*, vol. 37, no. 6, pp. 44–50, 2020. https://doi.org/10.1109/MS.2020.2998281